



# Arithmetic

## ICS312 Machine-Level and Systems Programming

Henri Casanova ([henric@hawaii.edu](mailto:henric@hawaii.edu))

# Addition and Subtraction

- Two instructions used for additions and subtractions: **add** and **sub**
- Both instructions can be used on a pair of signed numbers or on a pair of unsigned numbers
  - One of the big advantages of 2's complement storage
  - **No mixing of signed and unsigned numbers**
- **IMPORTANT:** The CPU does not know whether numbers stored in registers are signed or unsigned!
  - You, the programmer, must keep your own interpretation of the number consistent throughout your program
  - The CPU will happily add whatever registers together using binary addition
- These two instructions each may set some bits of the FLAG register:
  - The **carry** bit
  - The **overflow** bit
  - The **zero** bit (=1 if the result is equal to zero)
  - The **sign** bit (=1 if the result is negative)

# The Magic of 2's Complement

- I have two 1-byte values, A3 and 17, and I add them together:  
 $A3 + 17 = BA$
- If my interpretation of the numbers is **unsigned**:
  - $A3h = 163d$
  - $17h = 23d$
  - $BAh = 186d$
  - and indeed,  $163d + 23d = 186d$
- If my interpretation of the numbers is **signed**:
  - $A3h = -93d$
  - $17h = 23d$
  - $BAh = -70d$
  - and indeed,  $-93d + 23d = -70d$
- So, as long as I stick to my interpretation, the binary addition will do the right thing.... amazing!
  - Same thing for the subtraction

# Overflow???

- Generally speaking, overflow occurs when the result of an arithmetic operation generates a result that's “out of range”
- This happens because a register has a limited number of bits, which means that our interpretation of a number comes with a valid range
- For instance
  - adding 1-byte unsigned quantity 240d to 1-byte unsigned quantity 100d will lead to an overflow because  $340d > 255d$
  - subtracting 1-byte unsigned quantity 240d from 1-byte unsigned quantity 100d will lead to an overflow because  $-140d < 0d$
  - adding 1-byte signed quantity 100d to 1-byte signed quantity 120d will lead to an overflow because  $220d > 127d$
  - etc.
- Question: how do we detect overflow in a program?
  - Important otherwise we could be working with bogus numbers
- It depends on whether numbers are signed or unsigned...

# Overflow for Unsigned Operations

- There is an overflow with an unsigned operation (i.e., on unsigned quantities) if the carry bit is set
- If the carry bit is set, that means we'd need a larger quantity to hold the result
  - This also works for subtractions (instead of a carry, we have a "borrow", but it's still set in the carry bit)
- 1-byte Example (all in hex):
  - $FF + 02$                       Carry is set (result would be 101h)
    - $255 + 2 > 255$
  - $01 - 02$                       Carry is set (result cannot be negative)
    - $1 - 2 < 0$
  - $8A - 0F$                       Carry is not set (result is 7Bh)
    - $138 - 15 = 123$

# Overflow for Signed Operations

- There is an overflow with a signed operation (i.e., on signed quantities) if the overflow bit is set
  - This bit is set when the sign of the result does not agree with the signs of the operands
- 1-byte Example (all in hex, same as before):
  - FF + 02                      Overflow is not set (result is 01h)
    - -1 + 2 = +1
  - 01 - 02                      Overflow is not set (result is FFh)
    - 1 - 2 = -1
  - 8A - 0F                      Overflow is set (result would be < 80h)
    - 8A is negative, and is equal to -76h = -118d
    - -118 - 15 < -128, and thus cannot be represented as a 1-byte signed quantity

# Determining Overflow

- Another way to determine whether a particular signed operation would overflow is to look at the sign of the result and see if it makes sense
- Example: 1-byte operation  $8A + A2$ 
  - $8A$  is negative
  - $A2$  is negative
  - In hex  $8A + A2 = 2C$  (and a carry)
  - $2C$  is positive
  - The sum of two negative numbers should be negative, so we've experienced an overflow



# Overflow is your Responsibility

- The processor merely computes bits and puts them into the destination location as if everything were fine, and it's your responsibility to check the overflow!
- Let's look at two examples
  - An unsigned arithmetic example
  - A signed arithmetic example
- Note that we will see later how to “check” the Carry bit and the Overflow bit in the FLAGS register



# Unsigned Overflow

On web site as  
ics312\_overflow\_unsigned.asm

mov	al, 0F0h	; al = F0h
mov	bl, 0A3h	; bl = A3h
add	al, bl	; al = al + bl
movzx	eax, al	; increase size for printing
call	print_int	; print al as an integer

- As a programmer we decided to do some computation with **unsigned values**
- We put value F0h in al (unsigned F0h is decimal 240)
- We put value A3h in bl (unsigned A3h is decimal 163)
- We add them together
- The “true” result should be decimal  $240 + 163 = 403$ , which cannot be encoded on 8 bits (should be  $< 255$ )
- But the processor just goes ahead:  $F0 + A3 = 193h$ , and then drops the leftmost bits to truncate to a 1-byte value to get 93h!
- To call print\_int, we need the integer in eax, so we movzx al into eax
- print\_int print the decimal value corresponding to 00000093h, that is: 147!
- This is obviously wrong, and we can tell (or will be able to shortly) because the carry bit is in fact set to 1
- **Note that this is all correct if we assume signed values and replace movzx by movsx, but then our initial interpretation of the two values is different**

# Signed Overflow

On web site as  
ics312\_overflow\_signed.asm

mov	al, 09Ah	; al = 9Ah
mov	bl, 073h	; bl = 73h
sub	al, bl	; al = al - bl
movsx	eax, al	; increase size for printing
call	print_int	; print al as an integer

- As a programmer we decided to do some computation with **signed values**
- We put value 9Ah in al (signed 9Ah is decimal -102)
- We put value 73h in bl (signed 73h is decimal +115)
- We subtract bl from al
- The “true” result should be decimal  $-102 - 115 = -217$ , which cannot be encoded on 8 bits (should be  $\geq -128$ )
- But the processor just goes ahead:  $9A - 73 = 27h$
- To call print\_int, we need the integer in eax, so we movsx al into eax
- print\_int prints the decimal value corresponding to 00000027h, that is: 39!
- This is obviously wrong, and we can tell (or will be able to shortly) because the overflow bit is in fact set to 1
- **Note that this is all correct if we assume unsigned values and replace movsx by movzx, but then our initial interpretation of the two values is different**

# Multiplication

- There are two instructions to perform multiplications
- Multiplying **unsigned** numbers: **mul**
- Multiplying **signed** numbers: **imul**
- Why do we need two different instructions?
- Consider the multiplication of FF by FF
  - If we assume unsigned quantities, this is  $255 * 255 = 65035 = \text{FE0Bh}$
  - If we assume signed quantities, this is  $-1 * -1 = 1 = 0001\text{h}$

# The mul Instruction

- The size of the result of the multiplication is sometimes twice larger than the size of the operands
  - Multiplications just leads to much bigger numbers than additions
  - At most the result will be twice the size of the operands ( $255 * 255 = 65,025$ , which is encodable on 2 bytes)
- The oldest form of multiplication is the “mul” instruction, which produce a result twice the size of its **unsigned** operand

mul      <register or memory reference>

- If the operand is a byte, then it is multiplied by AL and the result is stored in (16-bit) AX
- If the operand is 16-bit, it is multiplied by AX and stored in (32-bit) DX:AX
  - There used to be no 32-bit registers
- If the operand is 32-bit, it is multiplied by EAX and the result is stored in (64-bit) EDX:EAX
  - We don't have 64-bit registers on a 32-bit architecture

# The imul instruction

- Imul, which is used for **signed** numbers, has three formats:

imul src

imul dst, src1

imul dst, src1, src2

- The different combinations are shown in Table 2.2 in the text book
- This table uses the typical way in which one specifies operands:
  - reg16: a 16-bit register
  - reg32: a 32-bit register
  - immed8: an 8-bit immediate operand (i.e., a number)
  - mem16: a word of memory
  - etc.
- Let's look at the table

# The imul instruction

Will not overflow  
(although the  
overflow bit may  
be set)

dst	src1	src2	action
	reg/mem8		AX = AL * src1
	reg/mem16		DX:AX = AX * src1
	reg/mem32		EDX:EAX = EAX * src1
reg16	reg/mem16		dst *= src1
reg32	reg/mem32		dst *= src1
reg16	immed8		dst *= immed8
reg32	immed8		dst *= immed8
reg16	immed16		dst *= immed16
reg32	immed32		dst *= immed32
reg16	reg/mem16	immed8	dst = src1*src2
reg32	reg/mem32	immed8	dst = src1*src2
reg16	reg/mem16	immed16	dst = src1*src2
reg32	reg/mem32	immed32	dst = src1*src2

# Division

- Two instructions:
  - **div** for unsigned quantities
  - **idiv** for signed quantities
- They perform **integer division**
  - e.g.,:  $19 / 4$  produces quotient = 4 remainder = 3
- Only one format for both:  
div/idiv src
- If src is an 8-bit quantity:
  - AX is divided by src
  - quotient stored into AL
  - remainder stored into AH
- If src is a 16-bit quantity:
  - DX:AX is divided by src
  - quotient stored into AX
  - remainder stored into DX

# Division

- If src is a 32-bit quantity:
  - EDX:EAX is divided by src
  - quotient stored into EAX
  - remainder stored into EDX
- Warning: it's very common for programmers to forget initializing DX or EDX before the division





# Negation

- There is a convenient instruction to negate an operand: **neg**
- It simply computes the 2's complement of a quantity
- Works on 8-bit, 16-bit, or 32-bit quantities
  - either in registers or in memory
- We'll ignore the content of Section 2.1.5 in the textbook



# Example Program in Textbook

- Section 2.1.4 shows a sample program that uses all the arithmetic operations we just saw
- There is nothing particularly difficult about it, especially because overflows are not handled (so the numbers entered had better be “small”)
- One interesting point: One cannot divide by an immediate value and must use a register
- Make sure you go through this example and understand how it works
  - You may want to run it as well



# Conclusion

- One has to be careful when doing arithmetic operations because the processor happily produces results but it's your responsibility to check for overflow/carry bits