

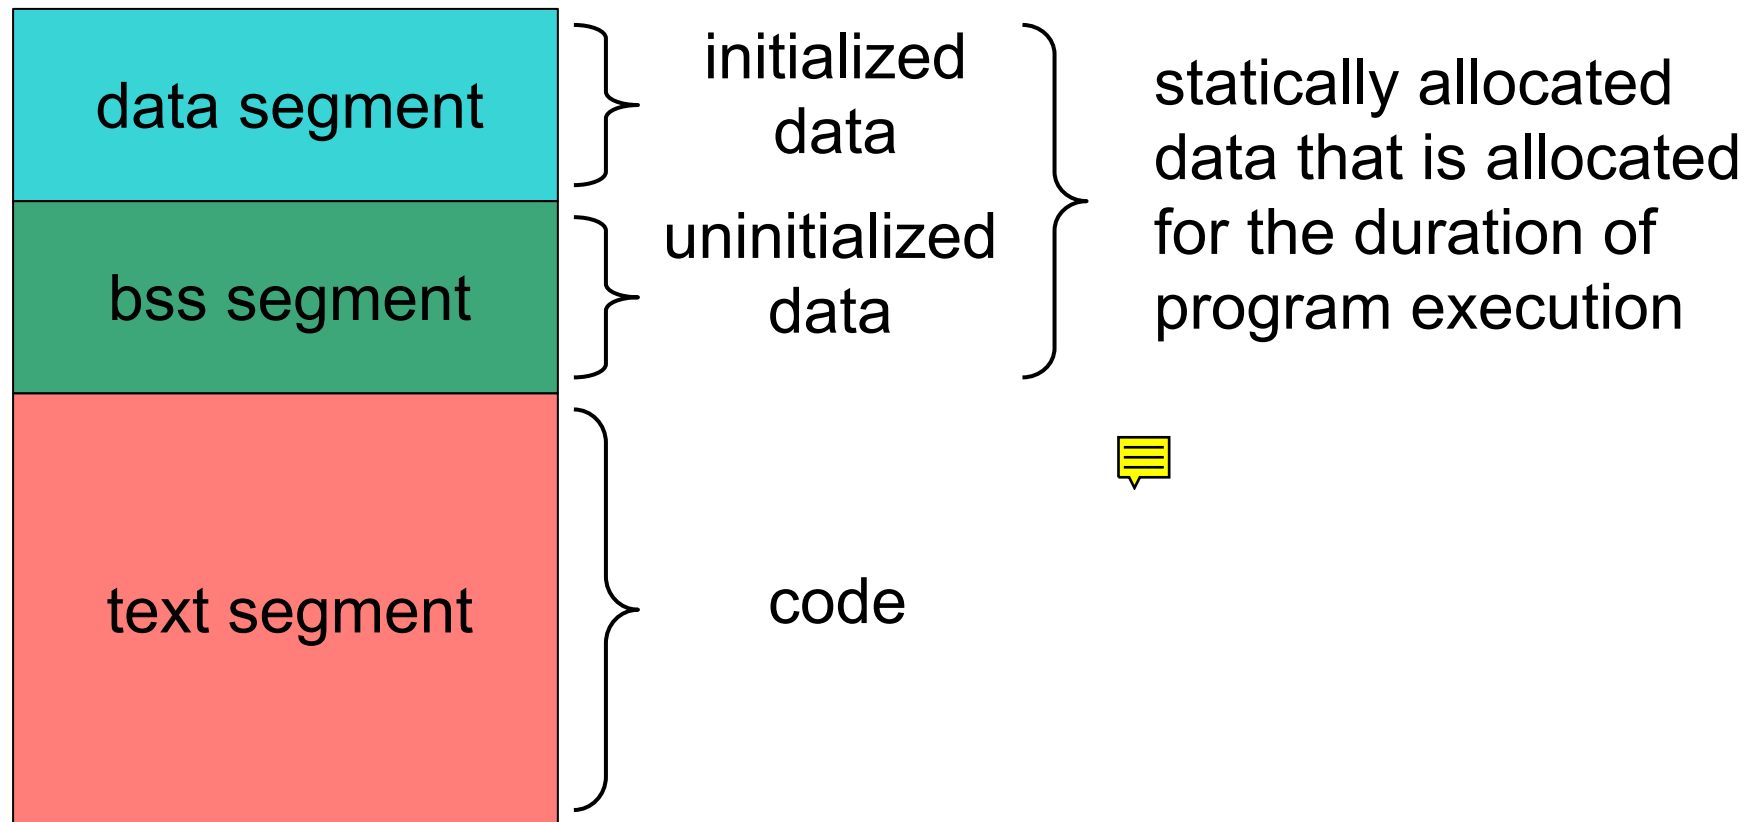


# **NASM: data and bss (inverted)**

## **ICS312 Machine-Level and Systems Programming**

Henri Casanova ([henric@hawaii.edu](mailto:henric@hawaii.edu))

# NASM Program Structure



# The data and bss segments

- Both segments contains **data directives** that **declare** pre-allocated zones of memory
- There are two kinds of data directives
  - **DX directives:** **initialized** data (D = “defined”)
  - **RESX directives:** **uninitialized** data (RES = “reserved”)
- The “X” above refers to the data size:

Unit	Letter(X)	Size in bytes
byte	B	1
word	W	2
double word	D	4
quad word	Q	8
ten bytes	T	10

# The DX data directives

- One declares a zone of initialized memory using three elements:
  - **Label**: the name used in the program to refer to that zone of memory
    - A pointer to the zone of memory, i.e., an address
  - **DX**, where X is the appropriate letter for the size of the data being declared
  - **Initial value**, with encoding information
    - default: decimal
    - b: binary
    - h: hexadecimal
    - o: octal
    - quoted: ASCII

# DX Examples

- L1 db 0
  - 1 byte, named L1, initialized to 0
- L2 dw 1000
  - 2-byte word, named L2, initialized to 1000
- L3 db 110101b
  - 1 byte, named L3, initialized to 110101 in binary
- L4 db 0A2h
  - 1 byte, named L4, initialized to A2 in hex (note the '0')
- L5 db 17o
  - 1 byte, named L5, initialized to 17 in octal ( $1*8+7=15$  in decimal)
- L6 dd 0FFFF1A92h (note the '0')
  - 4-byte double word, named L6, initialized to FFFF1A92 in hex
- L7 db "A"
  - 1 byte, named L7, initialized to the ASCII code for "A" (65d)



# ASCII Code

- Associates 1-byte numerical codes to characters
  - Unicode, proposed much later, uses 2 bytes and thus can encode  $2^8$  times more characters (room for all languages, Chinese, Japanese, accents, etc.)
- A few values to know:
  - 'A' is 65d, 'B' is 66d, etc.
  - 'a' is 97d, 'b' is 98d, etc.
  - ' ' is 32d

# DX for multiple elements

## ■ L8 db 0, 1, 2, 3

- Defines 4 bytes, initialized to 0, 1, 2 and 3
- L8 is a pointer to the first byte

## ■ L9 db "w", "o", 'r', 'd', 0

- Defines a **null-terminated** string, initialized to "word\0"
- L9 is a pointer to the beginning of the string

## ■ L10 db "word", 0

- Equivalent to the above, more convenient to write


# DX with the times qualifier

- Say you want to declare 100 bytes all initialized to 0
- NASM provides a nice shortcut to do this, the “times” qualifier
- `L11 times 100 db 0`
  - Equivalent to `L11 db 0,0,0,....,0` (100 times)



# Data segment example

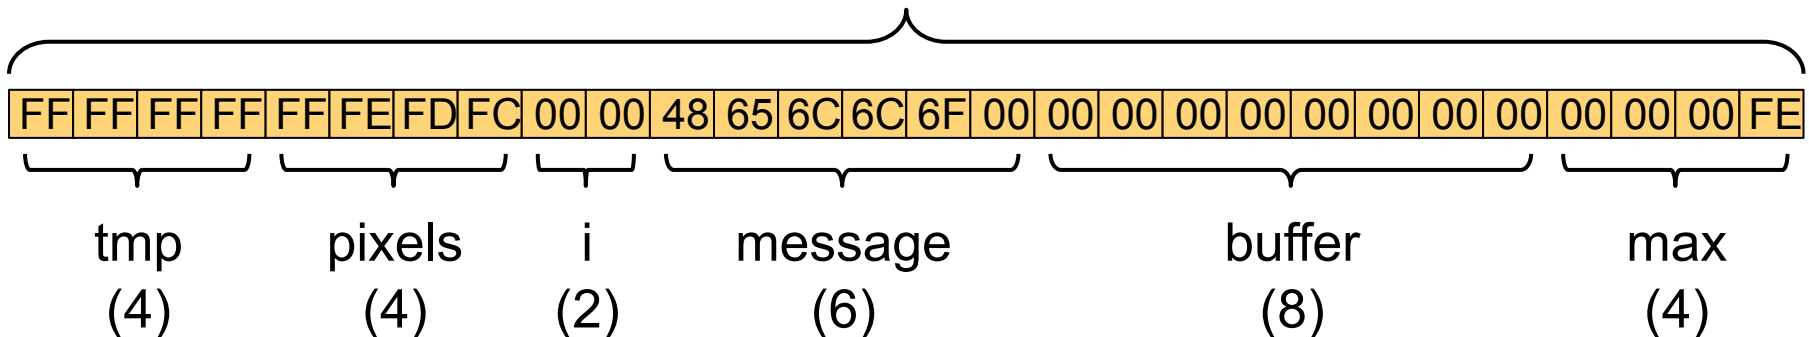
tmp	dd	-1
pixels	db	0FFh, 0FEh, 0FDh, 0FCh
i	dw	0
message	db	"H", "e", "llo", 0
buffer	times 8	db 0
max	dd	254




# Data segment example

tmp	dd	-1
pixels	db	0FFh, 0FEh, 0FDh, 0FCh
i	dw	0
message	db	"H", "e", "llo", 0
buffer	times 8	db 0
max	dd	254


28 bytes



# Endianness?



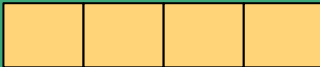
- In the previous slide we showed the above 4-byte **memory** content for a double-word that contains  $254 = 000000FEh$
- While this seems to make sense, it turns out that Intel processors do not do this!
  - Yes, the last 4 bytes shown in the previous slide are wrong
- The scheme shown above (i.e., bytes in memory follow the “natural” order): **Big Endian**
- Instead, Intel processors use **Little Endian**:



# Little Endian

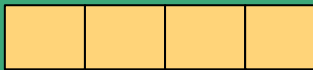
```
mov eax, 0AABBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

## Registers

eax 

ebx 

## Memory

[M1] 

# Little Endian

```
mov eax, 0AABBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

## Registers

eax AA BB CC DD

ebx

## Memory

[M1]

# Little Endian

```
mov eax, 0AABBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

## Registers

eax 

AA	BB	CC	DD
----	----	----	----

ebx 

--	--	--	--

## Memory

[M1] 

DD	CC	BB	AA
----	----	----	----

# Little Endian

```
mov eax, 0AABBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```

## Registers

eax AA BB CC DD

ebx AA BB CC DD

## Memory

[M1] DD CC BB AA

# Little Endian

```
mov eax, 0AABBCCDDh  
mov [M1], eax  
mov ebx, [M1]
```



In-register byte order and in-memory byte order, within a single multi-byte value, are different!



# Little/Big Endian

- Motorola and IBM processors use(d) Big Endian
- Intel/AMD uses Little Endian (used in this class)
- When writing code in a high-level language one rarely cares
  - Although in C one can definitely expose the Endianness of the computer
  - And thus one can write C code that's not portable between an IBM and an Intel!!!
- This only matters when writing **multi-byte** quantities to memory and reading them differently (e.g., byte per byte)
- When writing assembly code one often does not care, but we'll see several examples when it matters, so it's important to know this *inside out*
- Some processors are configurable (either in hardware or in software) to use either type of endianness (e.g., MIPS processor)


# Example

```
pixels      times 4      db      0FDh
x           dd      00010111001101100001010111010011b
blurb       db      "ad", "b", "h", 0
buffer      times 10     db      14o
min         dw      -19
```

- What is the layout and the content of the data memory segment on a Little Endian machine?
  - Byte per byte, in hex

# Example

pixels	times 4	db	0FDh
x	dd	00010111001101100001010111010011b	
blurb	db	"ad", "b", "h", 0	
buffer	times 10	db	14o
min	dw	-19	



# Uninitialized Data

- The RESX directive is very similar to the DX directive, but *always specifies* the number of memory elements
- L20        resw 100
  - 100 uninitialized 2-byte words
  - L20 is a pointer to the first word
- L21        resb 1
  - 1 uninitialized byte named L21



# Our first instructions

- At this point we need to introduce a few assembly instructions
  - adding integers
  - subtracting integers
  - moving data between registers / memory locations / constants

# Simple arithmetic and operands

- Assembly instructions can have operands, and it's important to know what kind of operands are possible
- **Register**: specifies one of the registers
  - **add eax, ebx**
  - means  $\text{eax} = \text{eax} + \text{ebx}$
- **Memory**: specifies an address in memory.
  - **add eax, [ebx]**
  - means  $\text{eax} = \text{eax} + \text{content of memory at address ebx}$
- **Immediate**: specifies a fixed value (i.e., a number)
  - **add eax, 2**
  - means  $\text{eax} = \text{eax} + 2$
- **Implied**: not actually encoded in the instruction
  - **inc eax**
  - means  $\text{eax} = \text{eax} + 1$

# Additions, subtractions

## ■ Additions

- `add eax, 4` ; `eax = eax + 4`
- `add al, ah` ; `al = al + ah`

## ■ Subtractions

- `sub bx, 10` ; `bx = bx - 10`
- `sub ebx, edi` ; `ebx = ebx - edi`

## ■ Increment, Decrement

- `inc ecx` ; `ecx++` (a 4-byte operation)
- `dec dl` ; `dl--` (a 1-byte operation)

# The move instruction

- This instruction moves data from one location to another  
`mov dest, src`
- Destination goes first, and the source goes second
- At most one of the operands can be a memory operand
  - `mov eax, [ebx]` ; OK
  - `mov [eax], ebx` ; OK
  - `mov [eax], [ebx]` ; NOT OK
- Both operands must be exactly the same size
  - For instance, AX cannot be stored into BL
- Examples:
  - `mov ax, ebx` ; NOT OK
  - `mov bx, ax` ; OK
- This type of “exceptions to the common case” make programming languages difficult to learn and assembly may be the worst offender
  - By contrast, Lisp is known for being very consistent (ICS313)



# Use of Labels

- It is important to constantly be aware that when using a label in a program, the label is a **pointer**, not a value
- Therefore, a common use of the label in the code is as a memory operand, in between square brackets '[' '']
- `mov AL, [L1]`
  - Move **the data at address L1** into register AL
- **Question:** how does the assembler know how many bits to move?
- Answer: it's up to the programmer to do the right thing, that is load into appropriately sized registers
- **Labels do not have a type!**
- **So although it's tempting to think of them as variables, they are much more limited: just pointers to a byte somewhere in memory**

# Moving to/from a register

- Say we have the following data segment

L      db      0F0h, 0F1h, 0F2h, 0F3h

- Example:            mov   AL, [L]

- AL: Lowest bits of AX, i.e., 1 byte
- Therefore, value F0 is moved into AL

- Example:            mov   [L], AX

- Moves 2 bytes into L, overwriting the first two bytes

- Example:            mov   [L], EAX

- Moves 4 bytes into L, overwriting all four bytes

- Example:            mov   AX, [L]

- AX: 2 bytes
- Therefore value F1F0 is moved into AX
- Note that this is **reversed** because of Little Endian!!

# More About Little Endian

- Consider the following data segment

L1      **db**      0AAh, 0BBh, 0CCh, 0DDh

L2      **dd**      0AABBCCDDh

- The instruction:      `mov    eax, [L1]`

puts DDCCBBAA into eax

- Note that we're loading 4x1 bytes as a 4-byte quantity

- The instruction:      `mov    eax, [L2]`

puts AABBCCDD into eax!!!

- Meaning that the memory content was DDCCBBAA

- When declaring a value in the data segment, that value is declared as it would be appearing in registers when loaded "whole"

- It would be confusing to write numbers in little endian in the program

- So all numerical values you write are in register-order not memory-order

# Example

- Data segment:

L1     db     0AAh, 0BBh

L2     dw     0CCDDh

L3     db     0EEh, 0FFh

- Program:

mov eax, [L2]

mov ax, [L3]

mov [L1], eax

- What's the memory content?


# Solution

## ■ Data segment:

L1     db     0AAh, 0BBh

L2     dw     0CCDDh

L3     db     0EEh, 0FFh



# Solution

L1		L2		L3	
AA	BB	DD	CC	EE	FF

mov eax, [L2] ; eax = FF EE CC DD

mov ax, [L3] ; eax = FF EE FF EE

mov [L1], eax ; L1 points to EE FF EE FF

L1		L2		L3	
EE	FF	EE	FF	EE	FF

Final memory content

# Moving immediate values

- Consider the instruction: `mov [L], 1`
- The assembler will give us an error: “operation size not specified”!
- This is because the assembler has no idea whether we mean for “1” to be 01h, 0001h, 00000001h, etc.
  - Labels have no type (they’re NOT variables)
- Therefore the assembler must provide us with a way to specify the size of immediate operands
- `mov dword [L], 1`
  - 4-byte double-word
- 5 size specifiers: byte, word, dword, qword, tword

# Size Specifier Examples

- `mov [L1], 1` ; Error
- `mov byte [L1], 1` ; 1 byte
- `mov word [L1], 1` ; 2 bytes
- `mov dword [L1], 1` ; 4 bytes
- `mov [L1], eax` ; 4 bytes
- `mov [L1], ax` ; 2 bytes
- `mov [L1], al` ; 1 byte
- `mov eax, [L1]` ; 4 bytes
- `mov ax, [L1]` ; 2 bytes
- `mov ax, 12` ; 2 bytes



# Brackets or no Brackets

- `mov eax, [L]`
  - Puts the content at address L into eax
  - Puts 32 bits of content, because eax is a 32-bit register
- `mov eax, L`
  - Puts the address L into eax
  - Puts the 32-bit address L into eax
- `mov ebx, [eax]`
  - Puts the content at address eax (= L) into ebx
- `inc eax`
  - Increase eax by one
- `mov ebx, [eax]`
  - Puts the content at address eax (= L + 1) into ebx

# Example

<code>first</code>	<code>db</code>	<code>00h, 04Fh, 012h, 0A4h</code>
<code>second</code>	<code>dw</code>	<code>165</code>
<code>third</code>	<code>db</code>	<code>"adf"</code>



<code>mov</code>	<code>eax, first</code>
<code>inc</code>	<code>eax</code>
<code>mov</code>	<code>ebx, [eax]</code>
<code>mov</code>	<code>[second], ebx</code>
<code>mov</code>	<code>byte [third], 110</code>

What is the content of “data” memory after the code executes on a Little Endian Machine?

# Example

```
first      db    00h, 04Fh, 012h, 0A4h
second     dw    165
third      db    "adf"
```

```
mov  eax, first
inc  eax
mov  ebx, [eax]
mov  [second], ebx
mov  byte [third], 110
```




# Example

first	db	00h, 04Fh, 012h, 0A4h
second	dw	165
third	db	"adf"

```
mov    eax, first
inc     eax
mov     ebx, [eax]
mov     [second], ebx
mov     byte [third], 110
```

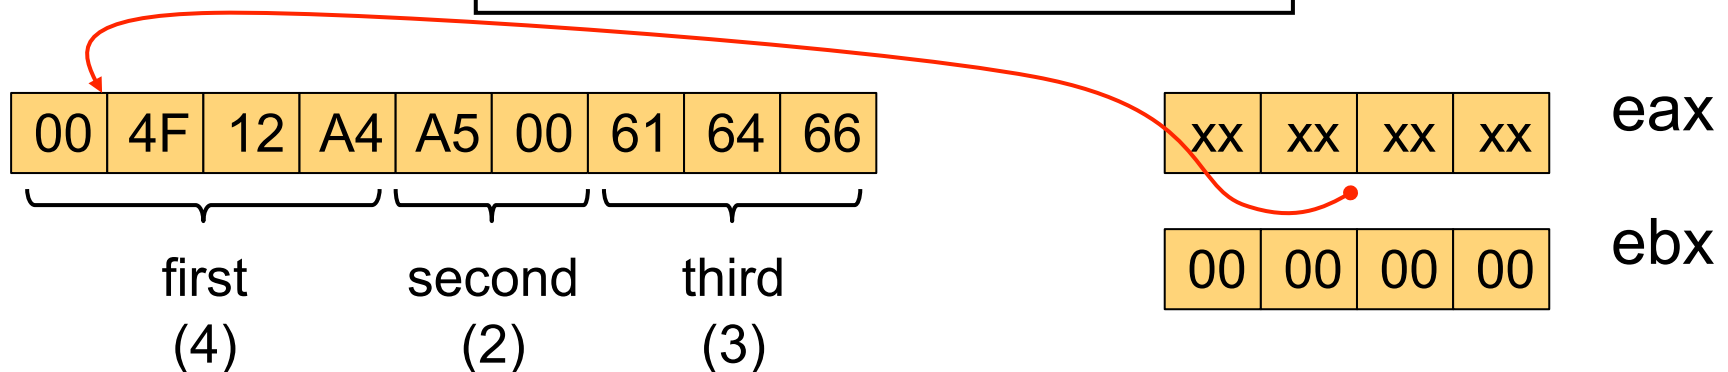
Put an **address** into **eax**  
(addresses are 32-bit)



# Example

first	db	00h, 04Fh, 012h, 0A4h
second	dw	165
third	db	"adf"

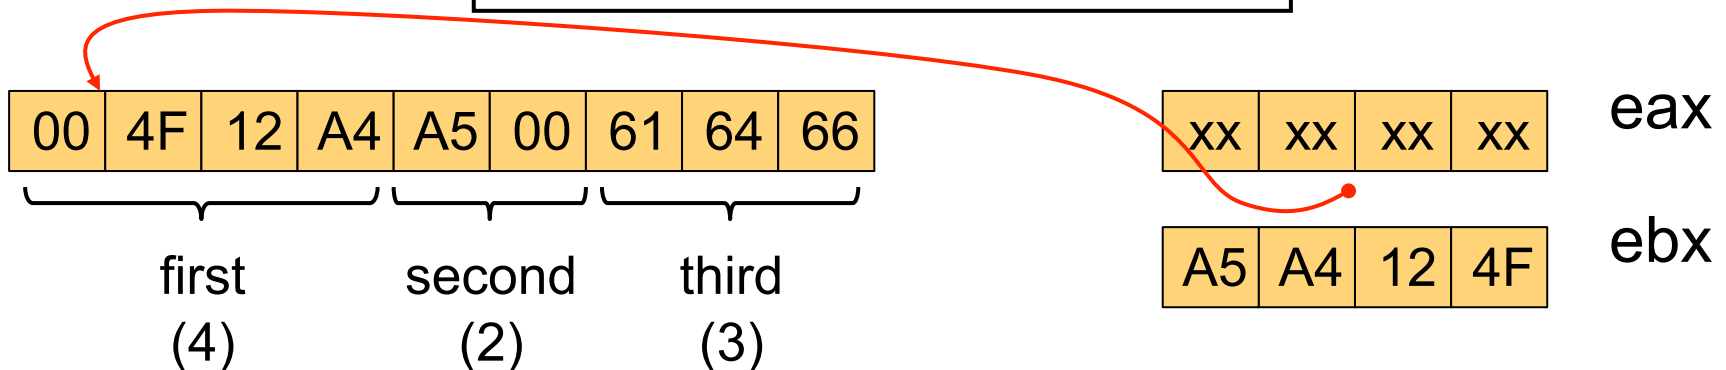
mov	eax, first
inc	eax
mov	ebx, [eax]
mov	[second], ebx
mov	byte [third], 11o



# Example

first	db	00h, 04Fh, 012h, 0A4h
second	dw	165
third	db	"adf"

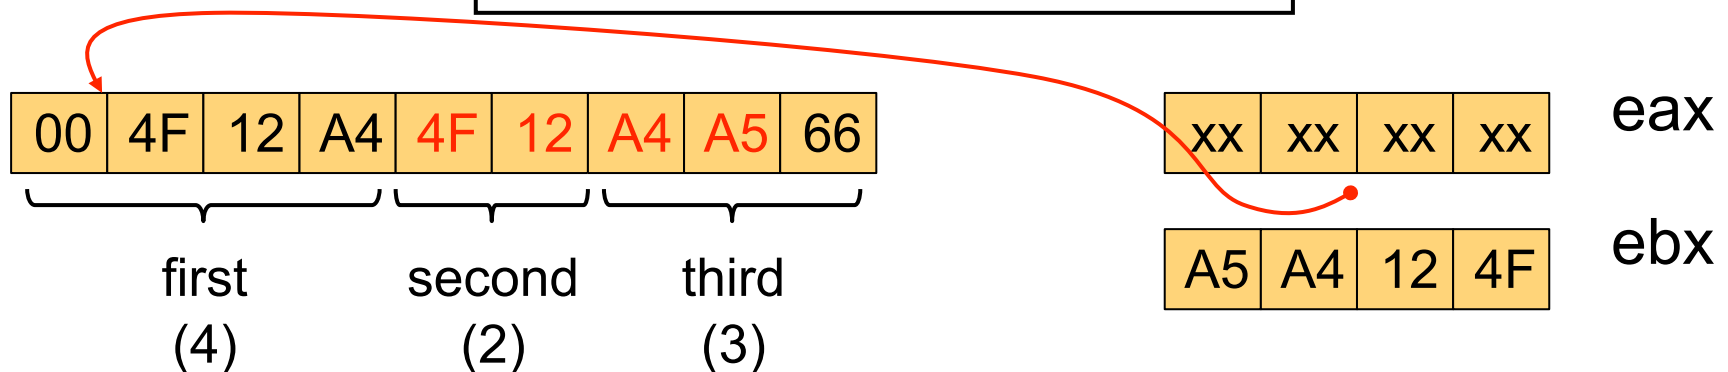
mov	eax, first
inc	eax
mov	ebx, [eax]
mov	[second], ebx
mov	byte [third], 11o



# Example

first	db	00h, 04Fh, 012h, 0A4h
second	dw	165
third	db	"adf"

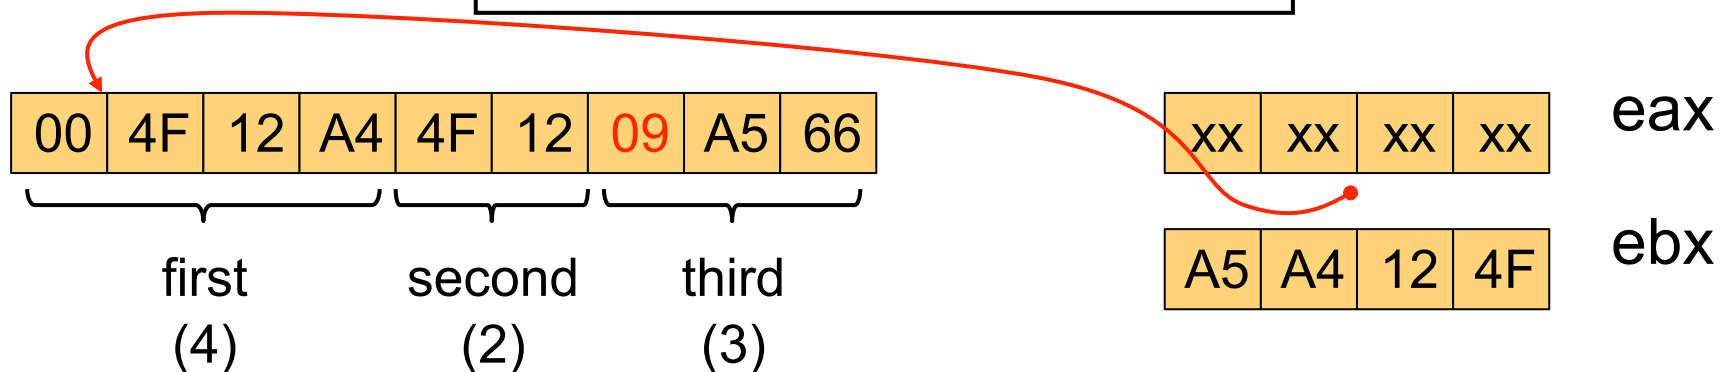
```
mov    eax, first
inc     eax
mov     ebx, [eax]
mov     [second], ebx
mov     byte [third], 110
```



# Example

first	db	00h, 04Fh, 012h, 0A4h
second	dw	165
third	db	"adf"

```
mov    eax, first  
inc    eax  
mov    ebx, [eax]  
mov    [second], ebx  
mov    byte [third], 11o
```








# Assembly is Dangerous

- The previous example is really a terrible program
- But it's a good demonstration of why the assembly programmer must be really careful
- For instance, we were able to store 4 bytes into a 2-byte label, thus overwriting the first 2 characters of a string that merely happened to be stored in memory next to that 2-byte label
- Playing such tricks can lead to very clever programs that do things that would be impossible (or very cumbersome) to do with many high-level programming language (e.g., in Java)
- But you really must know what you're doing
- Typically such behaviors are bugs

# x86 Assembly is Dangerous

- Another dangerous thing we did in our assembly program was the use of **unaligned memory accesses**
  - We stored a 4-byte quantity at some address
  - We incremented the address by 1
  - We read a 4-byte quantity from the incremented address!
  - This really removes all notion of a structured memory
- Some architectures only allow aligned accesses
  - Accessing an X-byte quantity can only be done for an address that's a multiple of X!





# Conclusion

- It's important to understand the memory layout