




Introduction to NASM Programming

**ICS312
Machine-Level and
Systems Programming**

Henri Casanova (henric@hawaii.edu)

Machine code

- Each type of CPU understands its own machine language
- Instructions are numbers that are stored in bytes in memory
- Each instruction has its unique numeric code, called the **opcode**
- Instruction of x86 processors vary in size
 - Some may be 1 byte, some may be 2 bytes, etc.
- Many instructions include operands as well



- Example:
 - On x86 there is an instruction to add the content of EAX to the content of EBX and to store the result back into EAX
 - This instruction is encoded (in hex) as: 03C3
- Clearly, this is not easy to read/remember

Assembly code

- An assembly language program is stored as text
- Each assembly instruction corresponds to exactly one machine instruction
 - Not true of high-level programming languages
 - E.g.: a function call in C corresponds to many, many machine instructions
- The instruction on the previous slides ($EAX = EAX + EBX$) is written simply as:

add eax, ebx

mnemonic

operands

Assembler

- An **assembler** translates assembly code into machine code
- Assembly code is NOT portable across architectures
 - Different ISAs, different assembly languages
- In this course we use the **Netwide Assembler (NASM)** assembler to write **32-bit Assembler**
 - See Homework #0 for getting NASM installed/running
- Note that different assemblers for the same processor may use slightly different syntaxes for the assembly code
 - The processor designers specify machine code, which must be adhered to 100%, but not assembly code syntax

Comments

- Before we learn any assembly, it's important to know how to insert comments into a source file
 - Uncommented code is a really bad idea
 - Uncommented assembly is a really, really bad idea
 - In fact, commenting assembly is **necessary**
- With NASM, comments are added after a ';'.
- Example:
 add eax, ebx ; y = y + b

Assembly directives

- Most assembler provides “directives”, to do things that are not part of the machine code but are convenient
- Defining immediate constants
 - Say your code always uses the number 100 for a specific thing, say the “size” of an array
 - You can just put this in the NASM code:
`%define SIZE 100`
 - Later on in your code you can do things like:
`mov eax, SIZE`
- Including files
 - `%include “some_file”`
- If you know the C preprocessor, these are the same ideas as
 - `#define SIZE 100` or `#include “stdio.h”`
- Use `%define` whenever possible to avoid “code duplication”
 - Because code duplication is evil



NASM Program Structure

; include directives

segment .data

; DX directives

segment .bss

; RESX directives

segment .text

; instructions

C Driver for Assembly code

- Creating a *whole* program in assembly requires a lot of work
 - e.g., set up all the segment registers correctly
- You will rarely write something in assembly from scratch, but rather only pieces of programs, with the rest of the programs written in higher-level languages like C
- In this class we will “call” our assembly code from C
 - The main C function is called a **driver**

```
int main()    // C driver
{
    int ret_status;
    ret_status = asm_main();
    return ret_status;
}
```

```
...
add eax, ebx
mov ebx, [edi]
...
```


So what's in the text segment?

- The text segment defines the `asm_main` symbol:

```
global    asm_main    ; makes the symbol visible
asm_main:                ; marks the beginning of asm_main
    ; all instructions go here
```

- On Windows, you need the `'_'` before `asm_main` although in C the call is simply to `"asm_main"` not to `"_asm_main"`
- On Linux you do not need the `'_'`
- I'll assume Linux from now on (e.g., in all the `.asm` files on the course's Web site)



NASM Program Structure

; include directives

segment .data

; DX directives

segment .bss

; RESX directives

segment .text

global asm_main

asm_main:

; instructions

More on the text segment

- Before and after running the instructions of your program there is a need for some “setup” and “cleanup”
- We’ll understand this later, but for now, let’s just accept the fact that your text segment will always looks like this:

```
enter  0,0
pusha
;
; Your program here
;
popa
mov    eax, 0
leave
ret
```



NASM Skeleton File

; include directives

segment .data

; DX directives

segment .bss

; RESX directives

segment .text

global asm_main

asm_main:

enter 0,0

pusha

; Your program here

popa

mov eax, 0

leave

ret



Our First Program

- Let's just write a program that adds two 4-byte integers and writes the result to memory
 - Yes, this is boring, but we have to start somewhere
- The two integers are initially in the .data segment, and the result will be written in the .bss segment

Our First Program

segment .data

```
integer1    dd    15    ; first int
integer2    dd    6     ; second int
```

segment .bss

```
result      resd    1    ; result
```

segment .text

global asm_main

asm_main:

enter 0,0

pusha

```
mov     eax, [integer1]    ; eax = int1
```

```
add     eax, [integer2]    ; eax = int1 + int2
```

```
mov     [result], eax      ; result = int1 + int2
```

popa

```
mov     eax, 0
```

leave

ret

File ics312_first_v0.asm
on the Web site



I/O?

- This is all well and good, but it's not very interesting if we can't "see" anything
- We would like to:
 - Be able to provide input to the program
 - Be able to get output from the program
- Also, debugging will be difficult, so it would be nice if we could tell the program to print out all register values, or to print out the content of some zones of memory
- Doing all this requires quite a bit of assembly code and requires techniques that we will not see for a while
- The author of our textbook provides a nice I/O package that we can just use, without understanding how it works for now

asm_io.asm and asm_io.inc

- The “PC Assembly Language” book comes with many add-ons and examples
 - Downloadable from the course’s Web site
- A very useful one is the I/O package, which comes as two files:
 - asm_io.asm (assembly code)
 - asm_io.inc (macro code)
- Simple to use:
 - Assemble asm_io.asm into asm_io.o
 - Put “%include asm_io.inc” at the top of your assembly code
 - Link everything together into an executable

Simple I/O

- Say we want to print the result integer in addition to having it stored in memory
- We can use the `print_int` “macro” provided in `asm_io.inc/asm`
- This macro prints the content of the `eax` register, interpreted as an integer
- We invoke `print_int` as:
 `call print_int`
- Let's modify our program

Our First Program

```
%include "asm_io.inc"

segment .data
    integer1      dd    15    ; first int
    integer2      dd     6    ; second int
segment .bss
    result        resd 1      ; result
segment .text
    global asm_main
asm_main:
    enter         0,0
    pusha
    mov           eax, [integer1]    ; eax = int1
    add           eax, [integer2]    ; eax = int1 + int2
    mov           [result], eax      ; result = int1 + int2
    call          print_int          ; print result
    popa
    mov           eax, 0
```

File ics312_first_v1.asm
on the Web site

How do we run the program?

- Now that we have written our program, say in file `ics312_first_v1.asm` using a text editor, we need to assemble it
- When we assemble a program we obtain an **object file** (a `.o` file)
- We use NASM to produce the `.o` file:

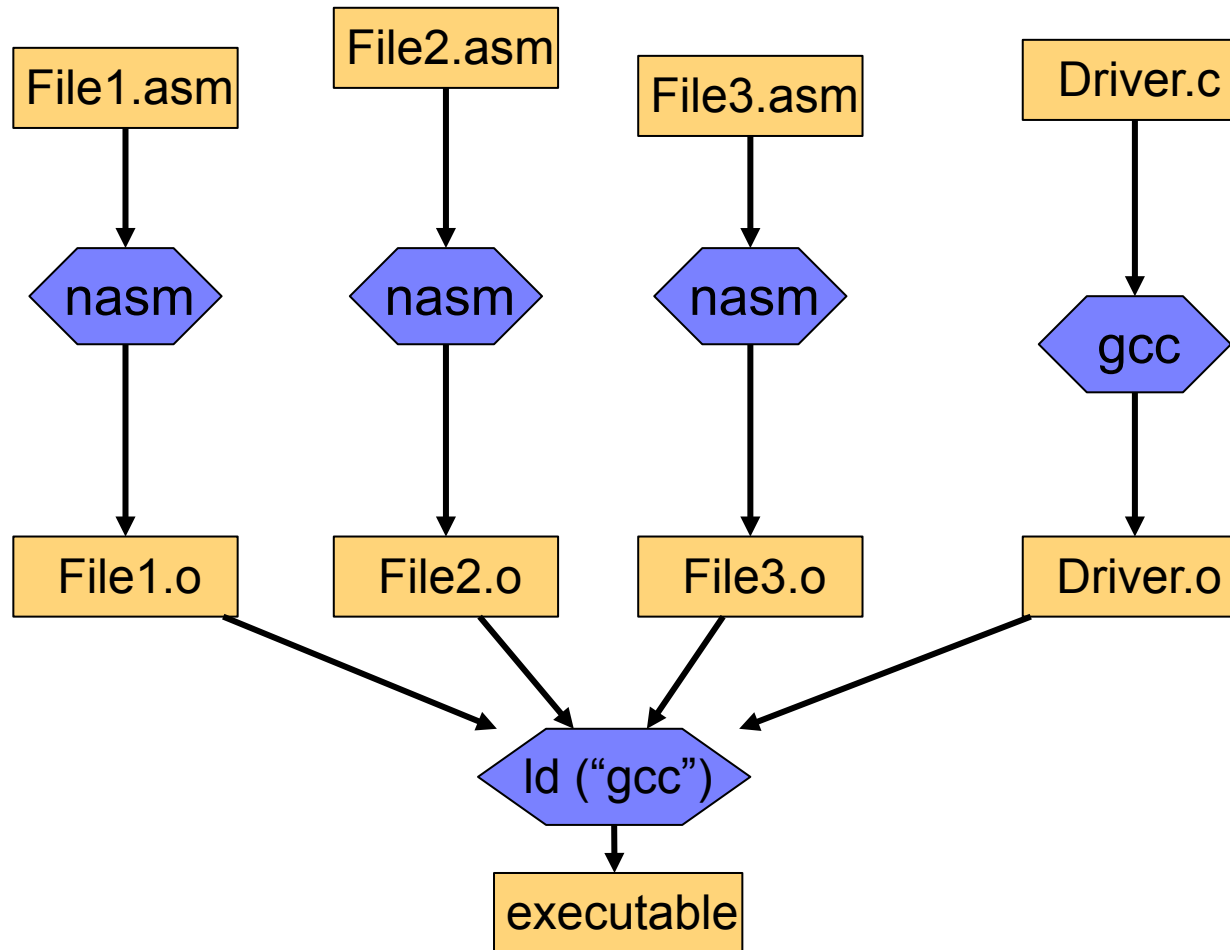
```
% nasm -f elf ics312_first_v1.asm -o ics312_first_v1.o
```
- So now we have a `.o` file, that is a machine code translation of our assembly code
- We also need a `.o` file for the C driver:

```
% gcc -m32 -c driver.c -o driver.o
```


 - We generate a 32-bit object (our machines are likely 64-bit)
- We also create `asm_io.o` by assembling `asm_io.asm`
- Now we have three `.o` files.
- We link them together to create an executable:

```
% gcc driver.o ics312_first_v1.o asm_io.o -o ics312_first_v1
```
- And voila... right?

The Big Picture



More I/O



- **print_char**: prints out the character corresponding to the ASCII code stored in AL
- **print_string**: prints out the content of the string stored at the address stored in eax
 - The string must be null-terminated (last byte = 00)
- **print_nl**: prints a new line
- **read_int**: reads an integer from the keyboard and stores it into eax
- **read_char**: reads a character from the keyboard and stores it into AL
- Let us modify our code so that the two input integers are read from the keyboard, so that there are more convenient messages printed to the screen

Our First Program

```
%include "asm_io.inc"
```

```
segment .data
```

```
msg1 db "Enter a number: ", 0
msg2 db "The sum of ", 0
msg3 db " and ", 0
msg4 db " is: ", 0
```

```
segment .bss
```

```
integer1 resd 1 ; first integer
integer2 resd 1 ; second integer
result resd 1 ; result
```

```
segment .text
```

```
global asm_main
```

```
asm_main:
```

```
enter 0,0
pusha
mov eax, msg1 ; note that this is a pointer!
call print_string
call read_int ; read the first integer
mov [integer1], eax ; store it in memory
mov eax, msg1 ; note that this is a pointer!
call print_string
call read_int ; read the second integer
mov [integer2], eax ; store it in memory
```

```
mov eax, [integer1] ; eax = first integer
add eax, [integer2] ; eax += second integer
mov [result], eax ; store the result
mov eax, msg2 ; note that this is a pointer
call print_string
mov eax, [integer1] ; note that this is a value
call print_int
mov eax, msg3 ; note that this is a pointer
call print_string
mov eax, [integer2] ; note that this is a value
call print_int
mov eax, msg4 ; note that this is a pointer
call print_string
mov eax, [result] ; note that this is a value
call print_int
call print_nl
popa
mov eax, 0
leave
ret
```

File ics312_first_v2.asm
on the Web site... let's compile/run it



Our First Program

- In the examples accompanying our textbook there is a very similar example of a first program (called first.asm)
- So, this is great, but what if we had a bug to track?
 - We will see that writing assembly code is very bug-prone
- It would be very cumbersome to rely on print statements to print out all registers, etc.
- So asm_io.inc/asm also provides two convenient macros for debugging!

dum_regs and dump_mem

- The macro dump_regs prints out the bytes stored in all the registers (in hex), as well as the bits in the FLAGS register (only if they are set to 1)

dump_regs 13

- '13' above is an arbitrary integer, that can be used to distinguish outputs from multiple calls to dump_regs

- The macro dump_memory prints out the bytes stored in memory (in hex). It takes three arguments:

- An arbitrary integer for output identification purposes
- The address at which memory should be displayed
- The number minus one of 16-byte segments that should be displayed
- for instance

dump_mem 29, integer1, 3

- prints out "29", and then $(3+1)*16$ bytes

Using `dump_regs` and `dump_mem`

- To demonstrate the usage of these two macros, let's just write a program that highlights the fact that the Intel x86 processors use Little Endian encoding
- We will do something ugly using 4 bytes
 - Store a 4-byte hex quantity that corresponds to the ASCII codes: "live"
 - "l" = 6Ch
 - "i" = 69h
 - "v" = 76h
 - "e" = 65h
 - Print that 4-byte quantity as a string

Little-Endian Exposed

```
%include "asm_io.inc"
```

```
segment .data
```

```
    bytes    dd      06C697665h ; "live"  
    end      db      0          ; null
```

```
segment .text
```

```
    global asm_main  
asm_main:  
    enter    0,0  
    pusha  
    mov     eax, bytes    ; note that this is an address  
    call    print_string  ; print the string at that address  
    call    print_nl      ; print a new line  
    mov     eax, [bytes]   ; load the 4-byte value into eax  
    dump_mem 0, bytes, 1   ; display the memory  
    dump_regs 0           ; display the registers  
    pusha  
    popa  
    mov     eax, 0  
    leave  
    ret
```

File
ics312_littleendian.asm
on the site...let's run it

Output of the program

The program prints
“evil” and not “live”

The address of “bytes”
is 0804A020”

“bytes” starts here

evil

Memory Dump # 0 Address = 0804A020

0804A020 65 76 69 6C 00 00 00 00 25 69 00 25 73 00 52 65 "evil????%i?%s?Re"

0804A030 67 69 73 74 65 72 20 44 75 6D 70 20 23 20 25 64 "gister Dump # %d"

Register Dump # 0

EAX = 6C697665 EBX = B7747FF4 ECX = BFBCB2C4 EDX = BFBCB254

ESI = 00000000 EDI = 00000000 EBP = BFBCB208 ESP = BFBCB1E8

EIP = 080484A4 FLAGS = 0282 SF

and yes, it's “evil”

The “dump” starts at
address 0804A020 (a
multiple of 16)

bytes in eax are
in the “live” order



Conclusion

- It is paramount for the assembly language programmer to understand the memory layout precisely
- We have seen the basics for creating an assembly language program, assembling it with NASM, linking it with a C driver, and running it